



Budapest University of Technology and Economics

Faculty of Electrical Engineering and Informatics

Department of Automation and Applied Informatics

Development of a template-driven website generator SaaS platform

MASTER'S THESIS

Author

Péter Hanich

Advisor

Gábor Knyihár

December 15, 2025

Contents

Abstract	i
1 Introduction	1
2 Literature review	2
2.1 Software as a Service	2
2.1.1 Aspects and challenges of SaaS	3
2.2 Laravel	4
2.2.1 Request Lifecycle	5
2.2.2 Ecosystem	6
2.3 React	6
2.3.1 SPA routing: TanStack Router	7
2.3.2 TanStack Query	7
2.3.3 Editor libraries	8
2.4 Kubernetes	8
2.4.1 Ingress, Gateway API	9
3 Architecture	10
4 Development	12
4.1 Backend	12
4.1.1 Setup	12
4.1.2 Eloquent ORM	12
4.1.3 Authentication	15
4.1.4 Authorization	15
4.1.5 Core Logic	16
4.1.6 Deployment Logic	17
4.1.7 Admin Panel	20
4.2 Frontend	22
4.2.1 Setup and overview	22

4.2.2	API client and context providers	24
4.2.3	Authentication, dashboard and account routes	25
4.2.4	Schemas, sites, and pages	29
4.2.5	Editor	29
4.3	Proxy application	29
5	Testing	30
5.1	Backend testing	30
5.2	Frontend testing	30
5.3	Usage example	30
6	Production considerations	31
7	Summary	32
	Acknowledgements	33
	Bibliography	34

HALLGATÓI NYILATKOZAT

Alulírott *Hanich Péter*, szigorló hallgató kijelentem, hogy ezt a diplomatervet meg nem engedett segítség nélkül, saját magam készítettem, csak a megadott forrásokat (szakirodalom, eszközök stb.) használtam fel. minden olyan részt, melyet szó szerint, vagy azonos értelemben, de átfogalmazva más forrásból átvettettem, egyértelműen, a forrás megadásával megjelöltem.

Hozzájárulok, hogy a jelen munkám alapadatait (szerző(k), cím, angol és magyar nyelvű tartalmi kivonat, készítés éve, konzulens(ek) neve) a BME VIK nyilvánosan hozzáférhető elektronikus formában, a munka teljes szövegét pedig az egyetem belső hálózatán keresztül (vagy autentikált felhasználók számára) közzétegye. Kijelentem, hogy a benyújtott munka és annak elektronikus verziója megegyezik. Dékáni engedéllyel titkosított diplomatervek esetén a dolgozat szövege csak 3 év eltelte után válik hozzáférhetővé.

Budapest, 2025. december 15.

Hanich Péter

Hanich Péter
hallgató

Abstract

Chapter 1

Introduction

Chapter 2

Literature review

2.1 Software as a Service

In order to develop a software as a service (SaaS) application suite we first have to explore what a SaaS entails. To do this an overview and understanding of cloud computing and its related terminologies is necessary. ISO defines cloud computing as a “paradigm for enabling network access to a scalable and elastic pool of shareable physical or virtual resources with self-service provisioning and administration on-demand [1].” The NIST (National Institute of Standards and Technology) defines five essential characteristics, four deployment models and three service models for the cloud [2].

The essential characteristics are the following:

- On demand self-service: resources can be provisioned without additional human assistance.
- Broad network access: users can access capabilities through standardized methods over the network.
- Resource pooling: the provider manages their cloud resources in such a way that it can serve multiple customers flexibly.
- Rapid elasticity: capabilities are scalable and re-assignable on demand.
- Measured service: the cloud monitors resource usage, and it is controlled and automatically optimized.

The aforementioned deployment models are:

- Private cloud: a single organizational entity uses the cloud.
- Community cloud: a community composed of multiple organizations or users uses the infrastructure.
- Public cloud: the cloud is open for use to the public.
- Hybrid cloud: the cloud is some mix of the models mentioned before.

All the above deployment models may be either on or off premise and additionally owned or managed by the organization(s) using them, third parties or some combination of the former.

Most importantly for the topic the three service models:

- Infrastructure as a Service (IaaS): the provided resources are fundamental to computing, such as network, storage, or processing power.
- Platform as a Service (PaaS): the provided resources allow the customer to deploy applications to the cloud providers platform without configuring the underlying computing infrastructure.
- Software as a Service (SaaS): “The capability provided to the consumer is to use the provider’s applications running on a cloud infrastructure. The applications are accessible from various client devices through either a thin client interface, such as a web browser (e.g., web-based email), or a program interface. The consumer does not manage or control the underlying cloud infrastructure including network, servers, operating systems, storage, or even individual application capabilities, with the possible exception of limited user-specific application configuration settings [2].”

There are many cloud providers available for use most notably Amazon Web Services (AWS), Microsoft’s Azure and Google Cloud Platform (GCP), which boast numerous amount services belonging to the models described previously.

2.1.1 Aspects and challenges of SaaS

After establishing a definition of SaaS and its related terminologies, lets dive into the different aspects and challenges of developing, maintaining and operating one.

“Software as a service (SaaS) is a software delivery model in which third-party providers provide software as a service rather than as a product over the Internet for multiple users [3].” SaaS applications hence are also web applications and built with web technologies and frameworks. However not all web applications are SaaS due to the aforementioned delivery model. In a way SaaS is a specific subset of web applications.

There is not a fixed defined set of characteristics that a SaaS has to meet in order to be considered one, however there are some commonly recurring ones that most take into account, such as multi-tenancy, customization and scalability [4].

Most SaaS applications support some form of multi tenancy. Multi tenancy is the ability to share the same (and single) hosted instance of the application between multiple tenants, instead of deploying the application separately for each tenant. A tenant is a user or group of users (e.g., an organization) with the same share of the service [5]. Commonly people refer to multi-tenancy as the ability to have multiple organizations use the service with their own share and users, however as we can see this is not a strict requirement. There are many ways to achieve multi-tenancy with different amounts of complexity in concern areas like architecture, scaling, and security. The methods usually differ by how separated each tenants share is on a database, data model and infrastructure level [6].

This goes hand-in hand with the next common characteristic which is customization. SaaS applications must support a level of customization to meet all tenants needs. This allows tenants to adapt the service more seamlessly into their workflows, processes, and requirements. The ability to change aspects also reduces repetitive tasks between users of the same tenants. Customization can range from interface level differences to full on tailoring of the tenants share of the application down to the platform or even infrastructure level. Each added aspect of personalization however increases complexity, due to which a balance must be struck between customization and standardization [7].

Scalability is the ability of the service to handle increasing workloads and usage needs. Generally there are two methods of scaling, vertical scaling (also known as scale-up) and horizontal scaling (also known as scale-out). Vertical scaling involves giving more resources to the machine running the service, while horizontal scaling means distributing the application on multiple machines. In a cloud environment consumers rely on a scale-out solution more by default due to the available amount of resources and also due to the fact that increasing a machine's resources is only feasible up to a limit. However, most complex production systems usually combine the two forms in some form. Not only do SaaS platforms commonly employ either a two-level or a generalized K-level scalability structure where each level is a separate dimension of the service capable of scaling independently. The platforms also make scaling solutions tenant-aware so that each tenant gets scaled for their respective needs [8].

Of course plenty more aspects could be reviewed and explored regarding SaaS applications, however the previously mentioned concerns and areas are sufficient to get a general understanding required for the topic.

Nowadays, many popular applications operate using a SaaS model such as Microsoft's Office suite, Slack, or Notion. In regard to the thesis topic there are also many similar existing solutions such as Squarespace, Wix, GoDaddy, or Shopify. However, they almost always require a subscription or one-time payment upfront and/or only offer a limited trial. They also frequently have a steep learning curve due to their complexity, and specialize in one given business area.

2.2 Laravel

In this section I will be reviewing Laravel, and its most important features regarding the thesis. Laravel is a popular PHP web framework created by Taylor Otwell and first released in June 2011 [9]. It is based upon Symfony which is another PHP web application framework first and foremost, but also contains a wide variety of PHP tools and libraries. It was also heavily inspired by Ruby on Rails.

"Laravel is, at its core, about equipping and enabling developers. Its goal is to provide clear, simple, and beautiful code and features that help developers quickly learn, start, and develop, and write code that's simple, clear, and lasting [9]." The tools and ecosystem of the framework allow developers to write more scalable and maintainable code.

Laravel uses the Model-View-Controller (MVC) architecture.

Models represent logical entities or resources in the application and is implemented by the Eloquent ORM (Object-Relational Mapping) of Laravel. Controllers provide methods to handle incoming requests and return an appropriate response or serve a view. Views display the aforementioned response from a controller method. Laravel provides the built-in Blade templating engine to return HTML views to the users [10].

An example diagram visualizing this architecture can be seen on Figure 1.

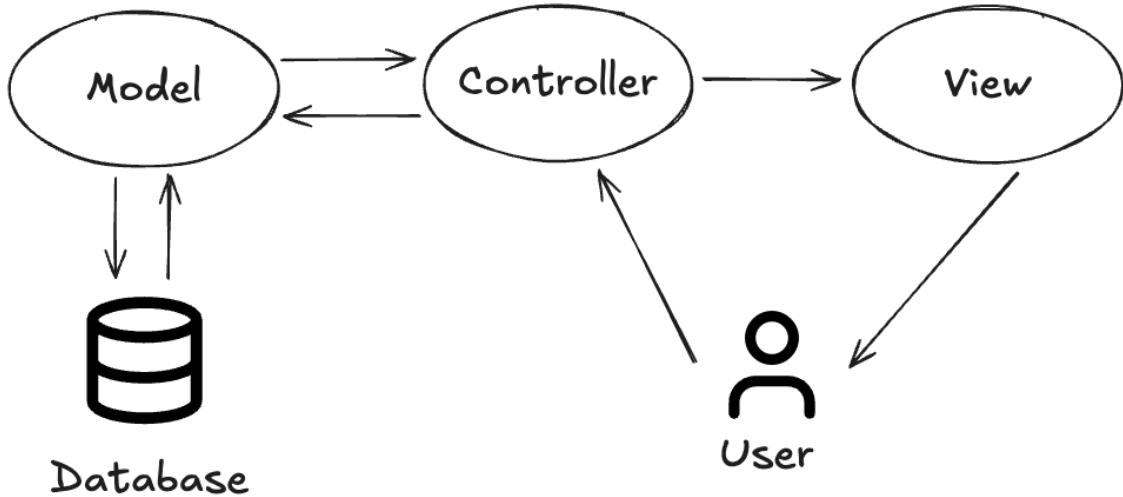


Figure 1: Model-View-Controller architecture

2.2.1 Request Lifecycle

Aside from understanding the general architecture of Laravel getting a grasp of its inner workings through the request lifecycle allows one to better understand on a high-level how the framework truly functions. As with all web applications a web server serves Laravel, most commonly Apache, Nginx, or nowadays FrankenPHP (built on Caddy). The starting point of the application is the public/index.php file to which the web server directs all requests. From here the file retrieves a Laravel application instance from the bootstrap/app.php file. Next it sends the incoming request to the HTTP kernel which bootstraps all necessary service providers, and executes application level middleware. These include the database, queue, validation, and routing components among many. Following this the router will dispatch the incoming request to route specific middleware, after which a route or controller will handle generating a response. Finally, the response will be sent out through the HTTP kernel and application instance [11].

A diagram of the complete request lifecycle and MVC can be seen on Figure 2

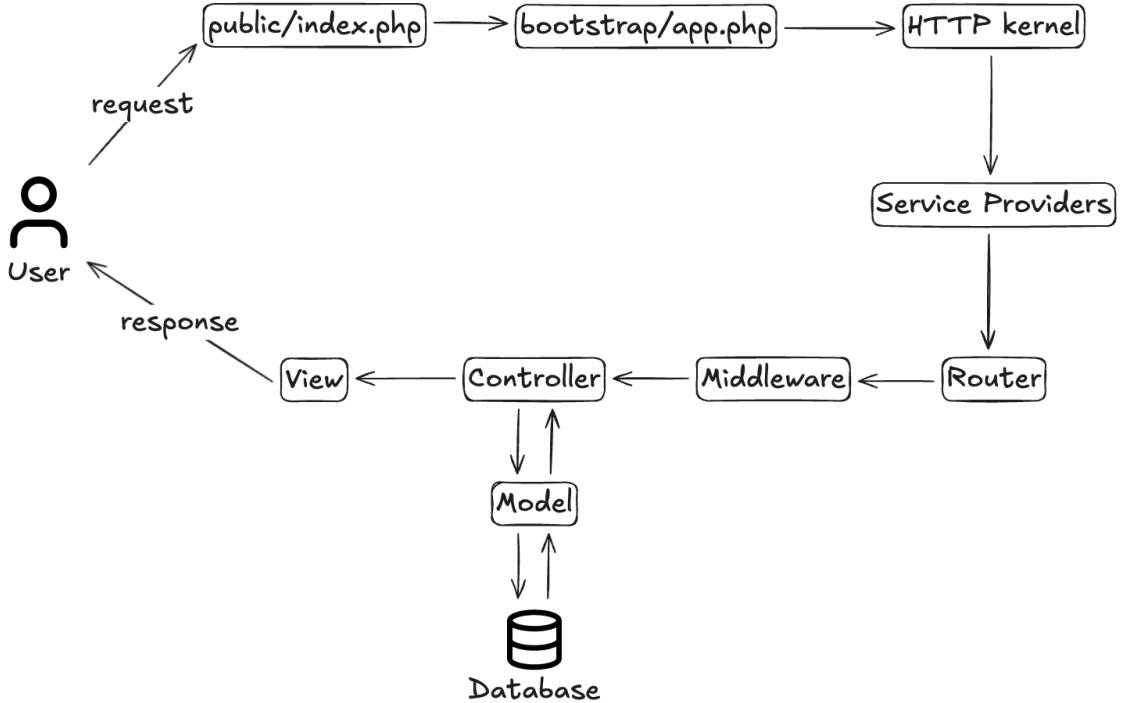


Figure 2: Laravel request lifecycle

2.2.2 Ecosystem

Laravel provides many built in functionalities for use, such as mailing, testing, queues, notifications, storage, authentication, and authorization among many. All of these capabilities can be extended, or new ones can be added through the rich package ecosystem surrounding the framework [12]. The creators of Laravel also maintain numerous first party packages, but there are also many other third party big packages and package maintainers. Most notably Spatie [13] whose team maintain many packages and tools for the framework and FilamentPHP, which is a plugin extendable admin panel and UI (user interface) framework [14].

2.3 React

This section provides a short overview of React and more importantly the packages I'm going to use throughout the implementation. React is JavaScript library for building user interfaces (UI) through reusable and composable components. It leverages a virtual DOM (Document Object Model) that decides whether the component has to be re-rendered or not based on the state and properties (props) of the component and their changes [15].

Developers commonly use React to implement client-side rendered single page applications (SPA), in which the server sends only one HTML file and the bundled JavaScript code dynamically rebuilds the DOM, without needing to make additional requests to the server [16].

2.3.1 SPA routing: TanStack Router

Since SPA-s only serve one HTML file from the backend they need some way to know which component to render for which web route. For this purpose developers use one of the many available routing libraries. In React the most popular routing library has been React Router for a long time [17]. However, recently React Router became part of the Remix framework which was later rebranded as version 7 of React Router and introduced declarative, data, and framework modes to the router [18]. Due to these changes and the fact that I planned to use TanStack Query (formerly known as React Query) I opted to use TanStack Router instead for easier use and better integration.

TanStack Router provides many features out of the box including but not limited to file and code based routing, caching, prefetching and data loading, route contexts, search, and URL parameter validation and native TypeScript support and type safe navigation [19].

The thesis project uses file-based routing, which we should take a closer look at. When we create our React application it creates a router instance with the default context (e.g., auth, theming) and settings like error boundaries, cache stale time and view transitions. We also have to create a routes folder and a root route where we can configure application wide error, not found and loading components. From here on out all tsx files (except the ones with special naming) will be treated as routes and will be auto managed by the router. This means creating route stubs and auto updating them on changes. From the structure of the routes' folder a route tree file will also be generated which provides the type safe route checking features. However, files can use special naming conventions for given purposes. I demonstrate the ones relevant for this project in Table 1.

Table 1: TanStack Router naming convention examples

Symbol	Example	Function
—	—root.tsx	The root route
— (prefix)	_layout.tsx	Layout route
- (prefix)	-button.tsx	Excluded, used for colocation
.	app.posts.tsx	Route nesting: /app/posts
\$	posts/\$postId.tsx	Route parameter

Folders can also be used for nesting instead of the dot notation with the prefixes outlined above. The main difference is in folder based grouping the file serving exact path of the folder (e.g., /posts) must be named index.tsx while layout routes should be named route.tsx. Folders can also be used for logical grouping without adding additional routes segments by putting the folder name in parentheses (e.g., (auth)) [19].

2.3.2 TanStack Query

“TanStack Query (formerly known as React Query) is often described as the missing data-fetching library for web applications, but in more technical terms, it makes fetching, caching, synchronizing and updating server state in your web applications a breeze [20].”

The library integrates seamlessly with TanStack Router and helps replace boilerplate API (Application Programming Interface) data fetching codes relying on fetch, useEffect and state management. Moreover, the package providers intelligent and automatic caching, background updates and request deduplication. TanStack Query has three fundamental

concepts, queries which fetch data (GET), mutations that modify data (PUT, POST, DELETE) and the query client which is the coordinator of the other two [21].

Under the hood the library uses Stale-While-Revalidate (SWR) caching which was put forward by RFC 5861 [22]. SWR is a client side caching strategy that works the following way. The client caches all queries upon fetching. The next time the user would need the result of that query the cache serves the **stale** data immediately, **while** triggering a re-fetch in the background. Once the server returns the fresh data it is **revalidated** and the application updates the UI without blocking. The client identifies queries by their query key which is most often the path of the query and all or any associated URL parameters. The cache can become stale after a set timeout, be manually invalidated, or on window or page switches and network reconnections [23]. This can be seen illustrated on Figure 3.

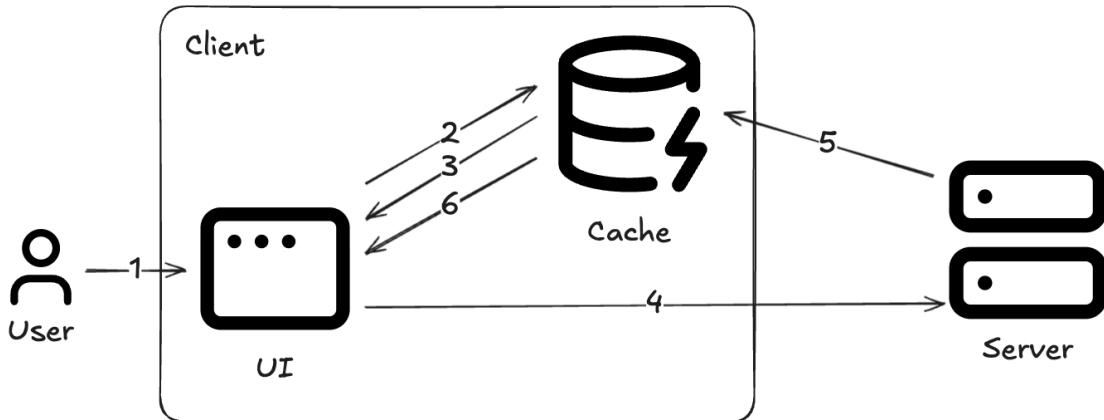


Figure 3: Stale-While-Revalidate caching

2.3.3 Editor libraries

Due to JavaScript's and by extension React's rich package and library ecosystem there exist multiple options to choose from regarding all needs. For the basis of the editor functionality I investigated two libraries, namely GrapesJS and CraftJS. They both offer capabilities to build one's own editor with drag and droppable UI elements [24]. Ultimately I decided against GrapesJS as the CraftJS is built more around React, while the former relies more heavily on pure JavaScript elements. CraftJS modularizes the editor and allows one to fully customize it according to their own needs [25].

2.4 Kubernetes

In this section I'm going to provide an overview of Kubernetes and its minimal subset of features required to understand this thesis, as Kubernetes is a vast and complex topic by itself. I'm assuming a basic understanding and familiarity of Docker, containerization, and virtualization from the reader.

“Kubernetes is an open source container orchestration system. It manages containerized applications across multiple hosts for deploying, monitoring, and scaling containers. Originally created by Google, in March of 2016 it was donated to the Cloud Native Computing Foundation (CNCF) [26].”

Kubernetes has multiple distributions, can be installed standalone, or used through the services of popular cloud platform providers. It is used declaratively to configure the desired state of system or application. Most often users interact with Kubernetes through its command line interface (CLI) with the `kubectl` command. This interacts with the underlying API of Kubernetes.

Pods are either a single or a group of containers with shared network and storage resources and the smallest configurable and manageable deployment unit in the system. A Pod can have any number of Init Containers. They must run successfully and in a set order, and can be used to do setup tasks for its respective Pod. Replicasets maintain a set number of replicas of the same pod while Deployments provide a declarative way to update Replicasets or Pods. Finally, a Service can be used to expose a network application that is running a set of pods behind a single outward facing endpoint [27].

2.4.1 Ingress, Gateway API

Since Pods can be short-lived, often recreated and have multiple replicas, a specialized service is needed to access HTTP routes from outside to cluster. This is provided by the Kubernetes Ingress or Gateway API.

“Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster. Traffic routing is controlled by rules defined on the Ingress resource [28].”

The Gateway API is in a way is a successor to the Kubernetes Ingress and provides a more modern and scalable solution to route and traffic management [29].

Both the Ingress and Gateway API need a controller or provider to function. One such industry standard is Traefik which can serve both functionalities, and is an open source application proxy [30].

A small illustration of the above concepts can be seen on Figure 4.

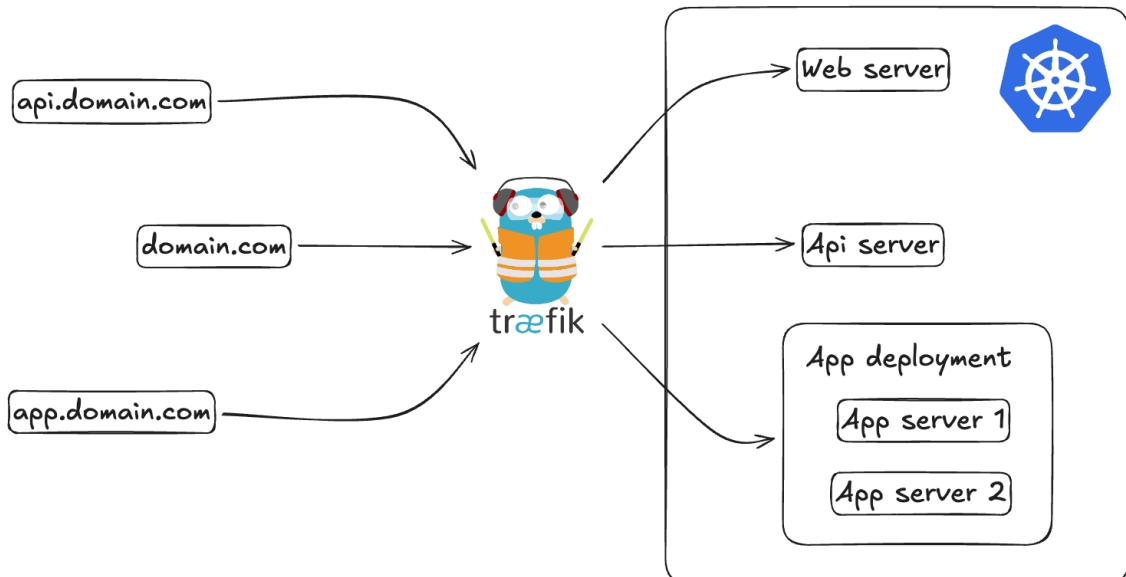


Figure 4: Traefik Gateway API

Chapter 3

Architecture

The goal of the thesis is to develop a template-driven website generator SaaS platform that allows users to create, customize, and publish single-page websites. Special attention must be paid to scalability, security, and user experience (UX). The system needs to provide users with the tools to create and design static websites for unique use cases, publish them and handle continuous integration and delivery of changes. All of this has to be implemented in a way that doesn't require users to have any knowledge of the underlying technical details. Thus, they should only have to be concerned about the content and appearance of their site.

A platform that can achieve this task from one end (the user) to the other (deployment) is inherently complex and requires the cooperation and integration of many technologies. To begin with, I will present the overarching high-level architecture of the complete system, introduce each component, then explain the ways they integrate and interact in detail. The architecture itself can be viewed on Figure 5.

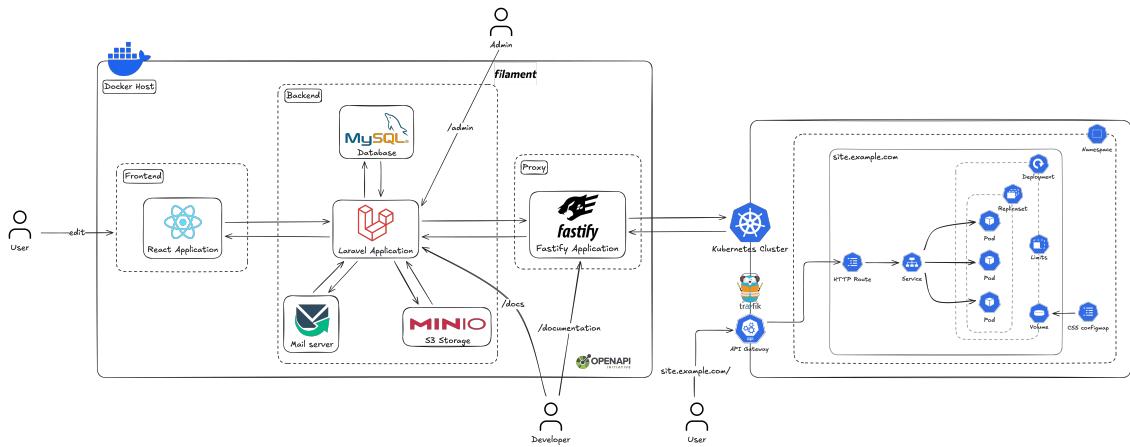


Figure 5: System level architecture

The system I designed has two main components a Docker host and a Kubernetes cluster. The Docker host contains the three main applications of the SaaS, each running inside containers and having their own Docker Compose configuration.

I built the frontend primarily using React with the Vite build and bundling tool. It communicates with the backend through a REST API. The CraftJS library provides the

drag and droppable editor capabilities which allow users to design and compose their websites using predefined blocks and templates.

The backend relies mainly on Laravel, but also uses MySQL for the database, Mailpit which is a mailing testing tool that acts as an SMTP server and MinIO, an S3 compatible object store. The backend serves the admin panel using FilamentPHP for the implementation, while I used the Scribe package to generate the backends API documentation using the Scalar documentation template. The backend communicates with both the frontend and the proxy application also through REST.

The Fastify NodeJS web framework provides the basis for the proxy application and the Kubernetes JavaScript library to connect to the running cluster. It also serves its own OpenAPI specification using the native Fastify OpenAPI plugin. For the reason of this functionality having its own separate application refer to Section 4.3.

The proxy connects programmatically to the Kubernetes cluster to deploy the created websites using its official JS client library. The websites are all deployed to a given namespace ('default' namespace if not changed) using a Kubernetes Deployment. The Deployments have the common style CSS (Cascading Style Sheets) attached as a volume through a ConfigMap. Furthermore, each Deployment by default has a 10 revision limit, default resource limits set, and a one replica requirement. In production these values could obviously be reconfigured and scaled up. Finally, the application also creates a Service for each website as well as an HTTP Route to connect to the Traefik API Gateway of the cluster for routing.

A user thus is able to create and design their website on the frontend which is then saved to the backend and through the proxy app deployed to the running Kubernetes cluster. Outside users can then visit the deployed website through the Traefik API Gateway acting as a reverse proxy.

Chapter 4

Development

4.1 Backend

4.1.1 Setup

I implemented the backend primarily using the Laravel PHP web framework as stated beforehand. As the first step the application has to be installed and set up. This can be done with Laravel's own installer, where many application templates can be selected. A React template is available, however that uses Inertia.JS by default which is a package that enables the frontend to be server side routed. I wished to opt out of this and develop a traditional API based SPA, so I decided to use the minimal setup. So Laravel will serve as an API backend and the view layer of the MVC architecture will be entirely handled by the frontend.

After installation, I decided to install the Laravel Sail package which provides a preconfigured development Docker Compose setup, and also provides options to install additional components into our environment. The extra containers in our case will be the MySQL database, the Mailpit SMTP test server and the MinIO S3 compatible object storage. However, the Laravel creators published Sail as a package which by default would need a local PHP and Composer (PHP's package manager) installation. This would kind of defeat its purpose. Luckily this can be avoided by using a temporary PHP container to quickly install initial dependencies and then use the Sail application for further required packages.

Afterwards the provided `.env.example` file needs to be copied to `.env` and the appropriate environment variables need to be set, including the mail server, database and object storage connections. In the MinIO object storage a bucket to store files must be created and set to public so that the links generated to the files can be accessed by the other applications. Finally, the application key must be initialized using the Artisan CLI and the default database migrations have to be run.

4.1.2 Eloquent ORM

Let me start building up the backend's feature set from the bottom, starting with the parts of the Eloquent ORM. The database schema of the backend including the most important tables can be seen on Figure 6.

Laravel has some other default tables, related to built-in features such as jobs and caching that I omitted from the diagram for better clarity as they will be unused in the application.

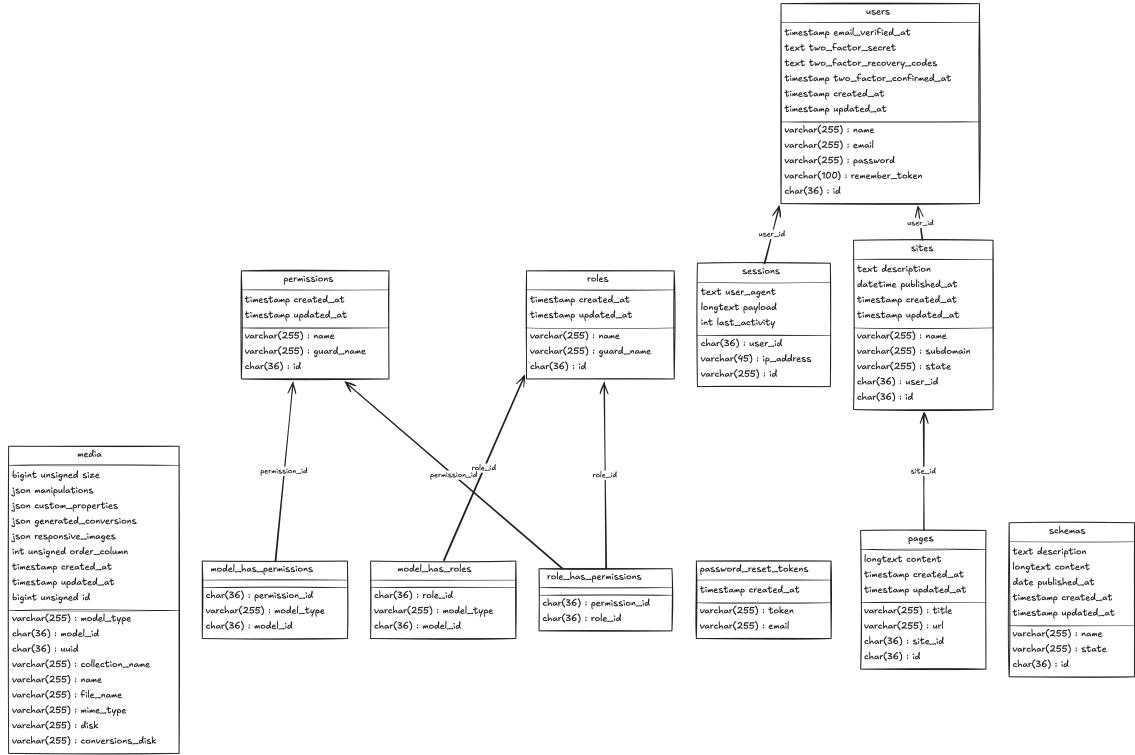


Figure 6: Database schema

All the tables have a corresponding migration file that defines a migration class with an up and down method. In these methods using the provided schema class we can define the blueprints of the tables with their fields and data types, and how to create and drop them. The Eloquent ORM then when running the migrations will take care of translating it to the query language of the chosen database management system.

The users table is auto generated, the only change I made is regarding the ID (Identifier) field. In all tables and models I opted to switch from the default integer ID to UUIDv7 (Universally Unique Identifiers version 7) which are time-based sortable universally unique identifiers composed of 32 hexadecimal digits. The application uses the simplest form of tenancy where each user is a tenant as that is sufficient for the current requirements.

All tables by default also have two timestamp fields a created_at and an updated_at. Laravel automatically manages these and casts them by default to the included carbon class, which provides helper methods to manage times professionally.

The Spatie Laravel Permission package adds the role related tables, which I will talk about in more detail in the authorization Subsection 4.1.4.

On the other hand the Spatie Media Library provides the media table which will be used to store profile pictures and the HTML files of the static pages. The application uses the password reset tokens and sessions table for authentication related features. The Sites table store the websites of the users. Each site must have a unique subdomain that will be assigned to it upon creation and which cannot be changed later.

Also, worth mentioning that each site also has a state that can be either draft or published, to track changes and whether they are visible to the end users. This is implemented using

the Spatie Model States library which adds an easy to configure state machine pattern support with helper functions such as checking viable transitions. It uses an abstract state (`PublishingState`) class with actual states being the concrete implementations (`Published`, `Draft`). Allowed transitions and transition handlers are also easy to configure and customize.

A site can also have pages which need a corresponding URL segment (e.g., `/about`) for routing. The database stores the pages content in a zlib compressed and base64 encoded format for better transfer and storage efficiency, but can be cast to the uncompressed version if needed. The structure of the content will be further specified in Section 4.2, but in short it is a JSON file encoding a tree structure. Pages also take advantage of the media library and S3 object storage to store the page's HTML in files. Last but not least the schemas table is a mix of the previous two, having both a content identical to that of the pages table and the same published and draft state as sites. Schemas allow administrators to create and publish page schemas that users can base their own pages off of.

The models paired with the tables contain most of the business logic that is independent of routes (heavy models). This includes defining all relationships and their inverses, like in the case of sites and pages a `hasMany`, and its inverse `belongsTo` functions, defining fillable fields that can be assigned from requests, and any custom casts and attributes. The former include casting to the aforementioned state classes while the latter includes accessing the decompressed and decoded content and the storage and retrieval of the HTML content of pages. Models register media collections for storing files by setting its name, accepted MIME (Multipurpose Internet Mail Extensions) types and the amount of files they can hold.

Finally, for most of the mentioned models I extended the special boot method with additional functionality. Inside this the CRUD functionality can be extended with additional logic. Examples include ensuring that sites transfer to the draft state when one of their pages gets updated, ensuring that a home page `(/)` exists for all sites, and relationship and media clean up logic upon model deletion.

A feature related to models is factories which allows the developer to create instances of models with set default attributes and persist them to the database. Another related feature is database seeding which allows us to populate the database with initial data. Since I use the aforementioned capabilities mostly for testing I will go into more detail about them in Section 5.1. For development purposes, however I use the user model's factory to create a super-admin user and seed roles and permissions into the database. I will talk about the previously mentioned features more in Subsection 4.1.4.

The final component related to the Eloquent ORM I used and want to mention here is API resource classes. They provide a way to transform models into the JSON format that is returned to the application's frontend. It can be thought of as a Data Transfer Object (DTO) implementation as resource classes allow fine granular control over the returned properties. Fields can be hidden from models, relationships can be included with select fields, and any custom or calculated field can be defined. Moreover, resource collection classes can be defined which define how to return a collection of the given model as opposed to returning a singular instance. Furthermore, resources classes can also be defined without them being connected to a given model.

For each model and endpoint I defined a corresponding resource and resource collection class. In order to parse request and responses from the proxy application I implemented custom deployment resources as well. I organized all of these classes according to the API versioning, to which refer to Subsection 4.1.5.

4.1.3 Authentication

Moving on from the ORM features, next I will describe the implementation of the authentication capabilities of the backend. The combination of two packages maintained by the Laravel team provide the functionality, the first being Laravel Sanctum, while the second is Laravel Fortify.

Laravel Sanctum provides API token based, SPA session based and Mobile token based authentication methods and provides CORS (Cross-Origin Resource Sharing) and CSRF (Cross-Site Request Forgery) protections for better security. Out of these three I only used the SPA session based authentication using cookies for the frontend. To begin with the SANCTUM_STATEFUL_DOMAINS environment variable has to be set with the frontend and proxy URLs. Next, in the application bootstrapper the stateful API global middleware has to be applied. Lastly the CORS configuration should be published, and the allowed paths have to be set.

In contrast, Laravel Fortify provides a frontend agnostic implementation of authentication features for the backend. In other words it implements the routes and controllers for login, registration, password reset and confirmation, email verification and two-factor authentication if needed. Each of these features can be enabled or disabled in its own configuration file. Additionally fortify can provide Blade templating based views for each of these features if needed however since I will use an SPA frontend I disabled them. I also opted-out from the two-factor authentication capabilities, the rest however are fully supported by the developed application suite.

One small caveat with using a separate React frontend is that the default password reset and email verification links that are generated by the backend navigate to itself instead of the frontend. I fixed this by overriding the link generator methods in the application service provider. To finalize the authentication implementation the user model has to be extended from the Authenticable class, and has to implement the MustVerifyEmail interface.

4.1.4 Authorization

Strongly related to authentication is authorization. Laravel provides gates and policies built-in for this purpose. Gates are closures that define if the application authorizes the given user to make a certain action.

However, I won't go into more detail about them as instead I will be using Spatie's Laravel Permission package. It allows developers to manage permissions and roles in a database. Since it provides native enum support I implemented a role and permission enum class.

The defined roles in the system are user, which grant basic permissions after registration, admin who can manage users, schemas and access the admin panel, and lastly the super admin. The last is the only role where I use Laravel's gates natively to grant super admins implicitly all permissions in the application service provider before any other checks.

I defined permissions separately for each model's and each CRUD method, as well as a unique permission for allowing access to the admin panel.

Since the package is storing authorization logic in the database it needs to be available after application setup. The straightforward method to do this is to write a seeder class that initializes the default values and what maps the appropriate permissions to the corresponding roles.

Now I only needed to set up the policy classes that use these roles. For each model an appropriate policy class can be created that checks whether the given user can be allowed to make the specified action. Inside these functions we can check users against ownership of the roles defined before. A special use case of policies I used is inside the page's policy where I check against the page's URL and prevent overriding or deleting the homepage.

4.1.5 Core Logic

The first part of the core CRUD logic is defining the API routes in the routes folder. Here routes can be defined for web, API, console use cases. The web handles standard web request if Laravel were to serve the frontend as well. Hence the React SPA takes care of this functionality the only route handler I defined here is a redirect to the frontend's index page. The console routes can be employed to define custom artisan commands, so it remained unused in this case.

I decided to use API versioning for a cleaner design and architecture, so I prefixed all of my routes with v1. If in the future a new API version would need to be released with breaking changes this allows for a cleaner transition while keeping backwards compatibility.

I grouped all routes by their respective controller that handles the incoming request. There is a controller for each model, an avatar controller for handling profile pictures, a dashboard controller for returning the frontend dashboard's content and a deployment controller for handling making deployment requests to the proxy application. This is worth talking about more in depth which will be done in Subsection 4.1.6. For each group the authentication I applied the middleware provided by Laravel Sanctum. Apart from the user and dashboard controller the others also have the verified middleware applied onto them that ensures the user making the request must have verified their email.

The routes also have the appropriate policy methods called on them that I talked about in Subsection 4.1.4, to check against the given user's roles and capabilities. Route parameters specified with e.g., {param} also get automatically dependency injected into the controller method.

Next in line I implemented Laravel's form request classes for the POST and PUT routes. They encapsulate the custom validation logic for each incoming request, which can be defined either using one of the many built-in validation rules or by creating a custom rule via a callback. A custom rule I implemented is checking that each path segment mapped to a page must be unique within the parent site. I used the predefined rules to validate against types, null checks, optional body parameters, sizes and file MIME types among many. Custom auth can be implemented here, however since it is already taken care by the route policy and middleware I only provided a simple check to make sure the user exists. I also defined a body parameters function in each form request to help the OpenAPI schema generation.

Due to the separation of concerns into the many capabilities most controllers are incredibly thin. Frequently they are only comprised of retrieving the validated form data, performing the CRUD operation, and returning the appropriate API resource instance. I used however the Scribe third-party package for generating an OpenAPI schema and hosted API documentation. The package auto discovers routes and controller methods and tries to deduce the parameters, bodies, and responses, but it can have errors or shortcomings. One way it can be helped is by giving the controller methods annotations or doc comments. I chose the latter and provided additional descriptions for the schema generation. It is also possible to document the routes and controllers provided by other packages, which in my

case was Laravel Fortify. This can be achieved by writing a custom YAML file describing those endpoints.

I also modified the configuration of the package so that the application publishes the generated schema on the /docs endpoint, which can be seen on Figure 7.

The screenshot shows the Scribe API documentation interface. On the left, there is a sidebar with a search bar and a tree view of API endpoints categorized by module: Authentication, Avatars, CSRF Protection, Dashboard, Deployments, Email Verification, Pages, and Postman. The 'Pages' section is expanded, showing 'Get Pages for a Site', 'Create Page for Site', 'Get Page for Site', and 'Update Page for Site'. The 'Create Page for Site' endpoint is selected and detailed on the right. The details include the method (POST), path parameters (site_id), body requirements (title, content, html), responses (201), and examples of the JSON payload and response object. A code editor on the right shows a JavaScript Fetch example for creating a page.

Figure 7: OpenAPI docs powered by Scribe

4.1.6 Deployment Logic

The deployment logic technically first begins at the page controller. When a page gets updated the frontend sends both the content and generated static HTML's body in an encoded manner for more efficient network transfer. The HTML body then gets base64 decoded and zlib decompressed on the backend to recover the raw HTML. Following this the HTML gets sanitized to remove any unknown tags and malicious scripts and ensure no security concerns got added by any third-party. I implemented this by writing an HTML sanitizer service. This service uses Symfony's `HTMLSanitizer` class that I preconfigured with the allowed and rejected elements and attributes. Then I registered this service as a singleton in the application provider, so that Laravel can automatically discover it and inject it into the Page controller constructor. As a part of the sanitization process if it is successful, the service afterwards injects the body into a Blade template. This template provides the root HTML and head tags and a link to the stylesheet, which will be later injected by the proxy application. In the end the complete HTML content gets returned and saved to the S3 storage and of course and API resource object gets sent to the frontend.

The MinIO S3 object storage showing the stored HTML files can be seen on Figure 8.

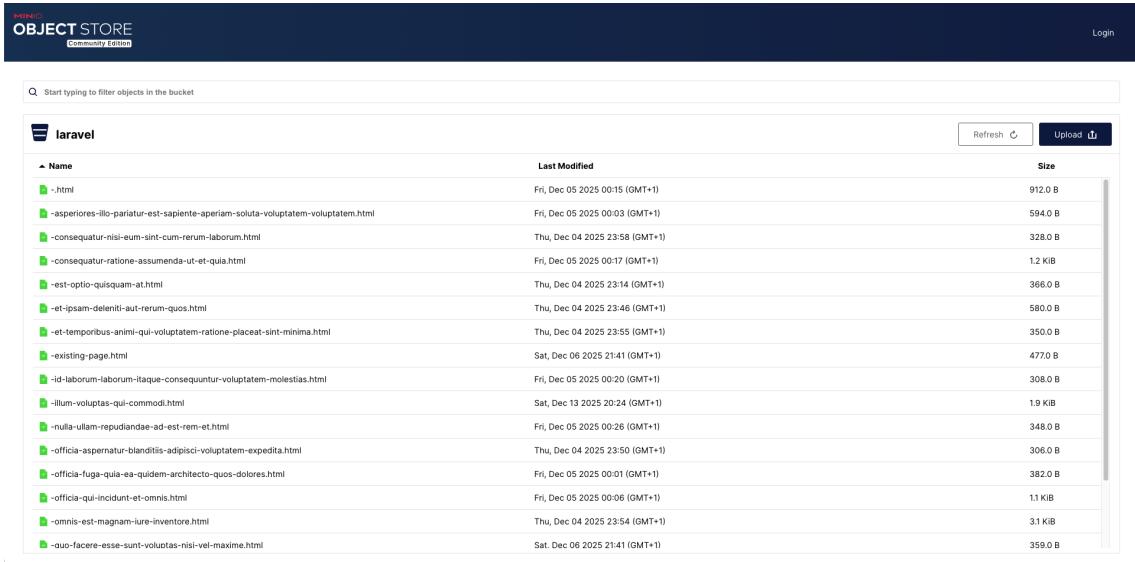


Figure 8: MinIO object storage

A sequence diagram displays this process to the fullest on Figure 9.

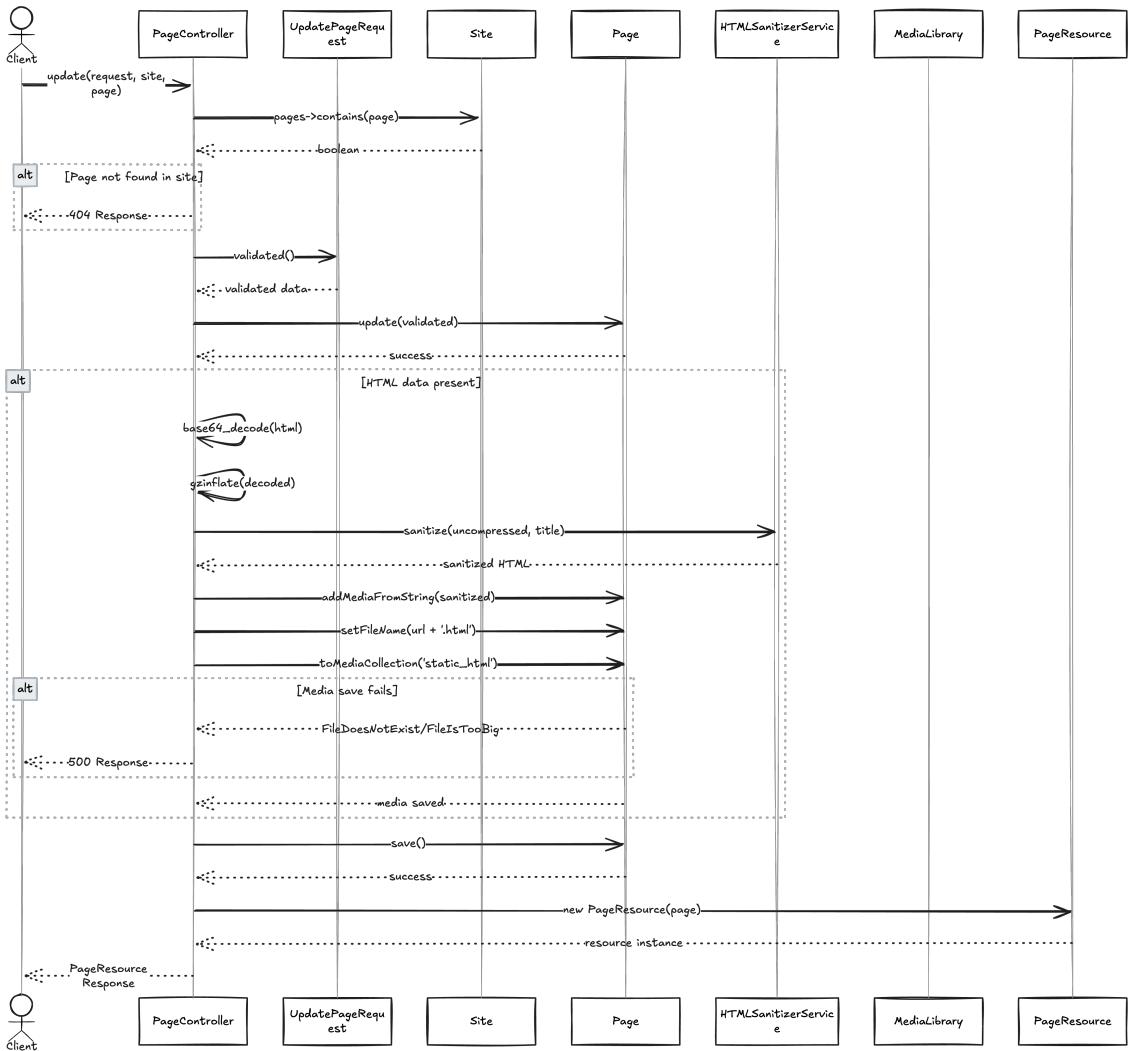


Figure 9: Sequence of saving the page content

Next comes the Deployment controller which handles the actual site publishing related methods, which includes a full CRUD and the option to restart the site if need be. I have written and registered another singleton service handling deployments that the controller forwards the incoming request to. The deployment service then checks if the site can be transferred to the published state, builds the path and body of the request which afterwards gets deferred to the proxy application. When the proxy returns with a response, the service forwards it to the controller which finally forwards it to the frontend.

Another sequence diagram on Figure 10 shows this process.

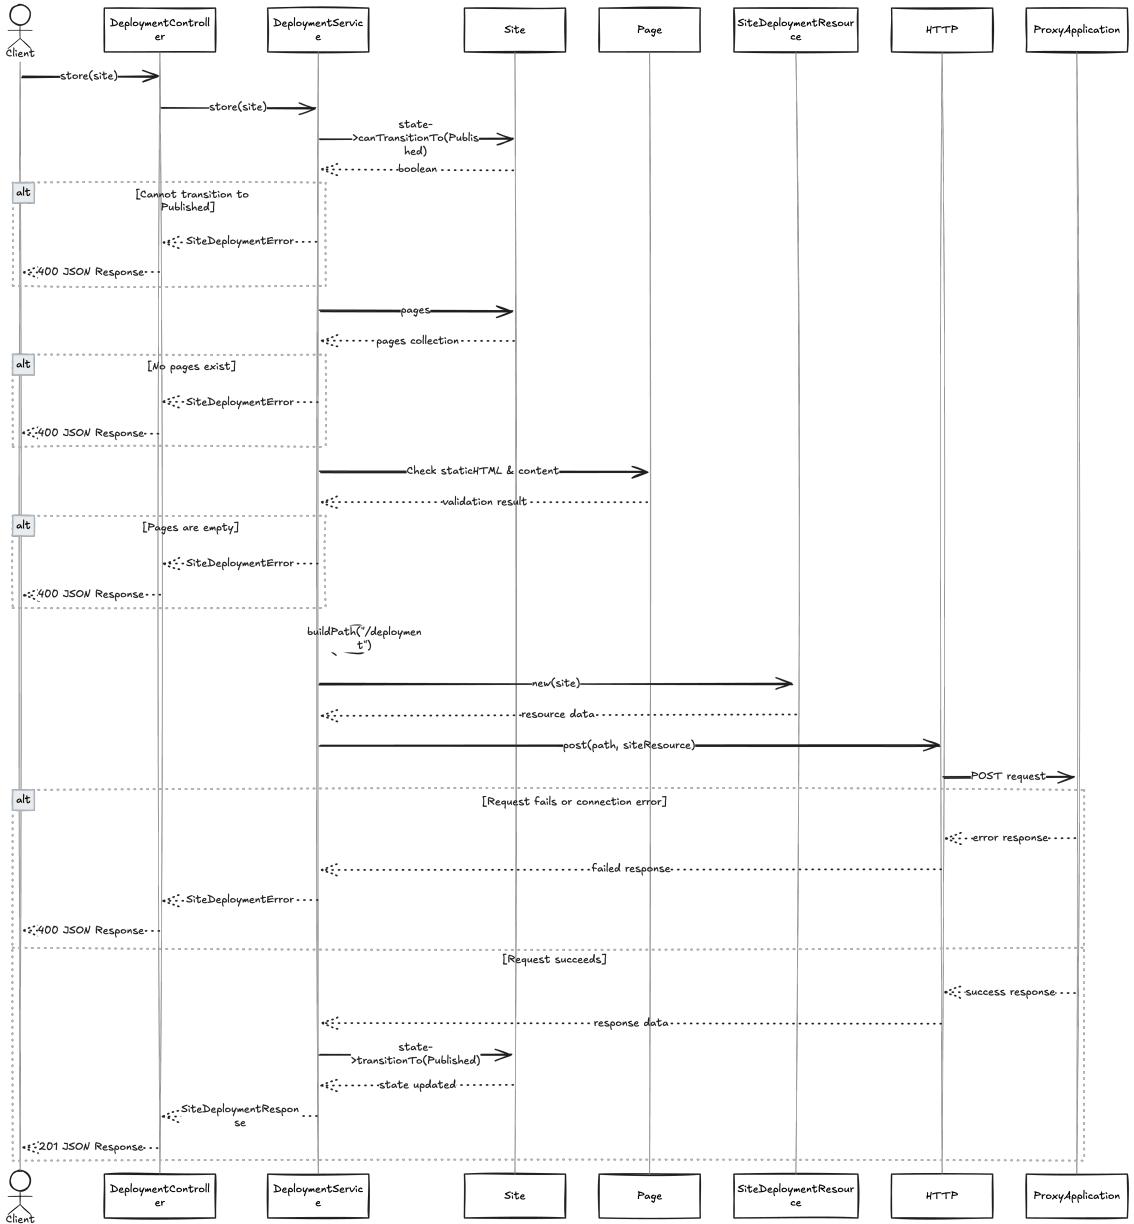


Figure 10: Sequence diagram of the deployment process

4.1.7 Admin Panel

The final feature of the backend I want to talk about in this chapter is the admin panel. It is provided by the Filament PHP package which specializes in building primarily admin panels for Laravel powered applications, but can also be used as a standalone server-driven UI framework if required. For the first step the user model has to implement the `FilamentUser` interface. This interface requires the implementation of the `canAccessPanel` method which has to return a boolean based on whether the user has access to the panel. Here is where I use the previously mentioned panel access permission. Afterwards using the included artisan command by filament the admin panel needs to be initialized. If done then it can be visited on the `/admin` route.

The admin site is empty at first and has to be populated. Filament defines resources which encompass the forms, pages, and schemas related to each Laravel model class. Here

UI elements Filament provides can be used to build interactive data tables, forms and views. The package provides filtering and sorting options by default, easily configurable searching and relationships can be displayed through more descriptive fields as well. For example instead of displaying the user's ID its name can be retrieved for a better UX. Further customization is possible, but the defaults cover most use cases.

On Figure 11 a data table view of sites can be seen.

The screenshot shows the 'Website Generator' admin panel with the title 'Sites' and a sub-title 'List'. On the left, there is a sidebar with 'Dashboard', 'Site Management' (Pages, Schemas, Sites), and 'User Management' (Users). The 'Sites' section is selected. The main area displays a table of site data with columns: ID, Name, Subdomain, and several status indicators (Draft, Published, etc.). A search bar and a filter modal are visible. The filter modal shows 'State' set to 'All' and an 'Apply filters' button. The table contains 10 rows of site data, each with a unique ID, name, subdomain, and various status and creation date details.

ID	Name	Subdomain	Draft	Published	Created
019b192b-ed46-73d7-bd75-14ed5f692b7e	Dare PLC	deckow			Dec 13, 2023
019b192b-ed51-70ca-a53d-b810bb53447d	Price Ltd	rodriguez			Dec 13, 2023
019b192b-ed55-73f0-be1d-9641ad0c1a56	Aufderhar Ltd	lynch			Dec 13, 2023
019b192b-ed5a-728a-933c-2f04128950d2	Mitchell, Crona and Oberbrunner	hansen			Dec 13, 2023
019b192b-ed61-70f3-9cd8-a7f56bbd3a32	Raynor PLC	little			Dec 13, 2023
019b192b-ed68-7152-aae5-77817f039acf	Carroll Inc	parisian			Dec 13, 2023
019b192b-ed70-7368-a9f1-36815a2581f7	Stracke Ltd	breitenberg			Dec 13, 2023
019b192b-ed76-73a3-8a53-f12235939ea2	Olson, Conn and Wilkinson	mcclure			Dec 13, 2023
019b192b-ed7e-7019-a3bf-5a5b54e2d5b8	Smitham-Shields	nienow			Dec 13, 2023

Figure 11: List view of sites on the admin panel

I extended the functionality of the admin panel with the Filament StateFusion plugin created by Assem Alwaseai and the Spatie Media Library plugin maintained by the Filament team themselves. The former plugin integrates the Spatie Model States package into the admin panel and provides quick actions to transfer models from one state to another. It also checks for allowed states too. The latter plugin provides a droppable file upload field for storing files in the media library's collections and also handles displaying them. This is mainly used to handle the user avatars from the panel.

Furthermore, I implemented a Filament relation manager for sites and pages so that the pages of the sites could be controlled from the site view. Lastly I created a custom action group to handle deployment related operations from the admin panel. These directly call the deployment service from the admin panel.

An overview of the previous two functionalities of a site's detail view can be seen on Figure 12.

The screenshot shows the 'Website Generator' admin interface. In the top left, there's a sidebar with 'Dashboard', 'Site Management' (Pages, Schemas, Sites), and 'User Management' (Users). The main content area is titled 'View Mitchell, Crona and Oberbrunner'. It displays site details: ID (019b192b-ed5a-728a-933c-2f04128950d2), Subdomain (hansen), Name (Mitchell, Crona and Oberbrunner), State (Draft), and a description (Totam officia repellat labore molestias necessitatibus officiis qui iure.). Below this, it shows user information: User (Adrianna Will), Created at (Dec 13, 2025 19:24:27), Updated at (Dec 13, 2025 19:24:27), and Published at (May 26, 2025 16:06:52). A 'Pages' section shows a table with one result: ID (019b192b-ed5c-73d3-aa67-00b45915dbe9), Title (Home), Url (/), Site (Mitchell, Crona and Oberbrunner), and a 'View' link. On the right, there are buttons for 'Edit', 'Deployment options' (with 'Update Deployment', 'Delete Deployment', and 'Restart deployment'), and a search bar.

Figure 12: Site detail view on the admin panel

4.2 Frontend

4.2.1 Setup and overview

I implemented the frontend using React first and foremost. The most important packages I used are TanStack Router for client side routing, TanStack query for robust data fetching and caching, and CraftJS for the editor capabilities.

As the first step, the React app has to be installed. I scaffolded the application using the Vite bundling tool's CLI and chose a TypeScript based application. Then I set TypeScript to use strict type checking to disallow any types, and also installed TailwindCSS and its Vite plugin as I used the ShadCN UI components. ShadCN distributes its components and blocks in an open code manner through the npx CLI, which means that instead of installing a npm package, the source code of the block gets copied into the application. This has the advantage that the components can be completely customized if needed, and the downside that the repository size grows considerably.

I also used predesigned UI blocks from ShadCN.IO [31]. While I was building the application all components from this site were open-sourced and free, however as of writing this documentation they introduced many new blocks that got locked behind a subscription. Previously free blocks remained as such, but require a free account now. This change doesn't affect the application as the source code got inserted into the project directly so further use and modification is possible.

Next I used Docker's init command to scaffold a Dockerfile and docker-compose.yaml for containerization, to which I only made minor changes, mostly regarding setting up environment variables. I also configured ESLint, which is a static code analyser, to ensure better code quality.

Since unlike Laravel, React is unopinionated about the structure of the application, let me talk about how I architected the frontend's folder structure. I was partly inspired by the

Bulletproof React repository [32], however I made a few changes to better fit my needs. The high level overview of the layout can be seen on Figure 13.

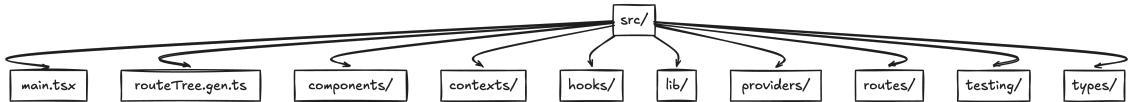


Figure 13: React application structure

The `main.tsx` is the setup file for the frontend application, about which I will go into detail shortly. The `routeTree.gen.ts` contains the routing tree that TanStack router generates. For more information about this refer to Subsection 2.3.1. The `components` folder contains the UI elements added by the ShadCN library and other component blocks I created using them that are global to the application. The context includes the globally defined application contexts, that I will talk about in Subsection 4.2.2. Inside `hooks` live the custom React hooks that I defined. `Lib` collects utility functions, global constants and the generated API client, which I will elaborate in Subsection 4.2.2. The `providers` folder houses the context providers that the aforementioned context objects return. Routes encompass not only the route hierarchy but using the router's - notation it also collocates any route specific components and logic. The `testing` folder accommodates all test related files. Refer to Subsection 5.2 for a thorough explanation. Finally, `types` include all TypeScript types that are either generated by the API client or Zod schemas.

Zod is a library used for schema validation primarily focusing on TypeScript. I used it in conjunction with React Hook Form, a hook based form library for React, to generate form schemas and provide client side validation. I also validated search parameters using it in a few routes. A great feature of Zod is that TypeScript types can be inferred from form schemas for further use.

Now let me talk about setting up the application in the `main.tsx` file. Some details will be expanded upon later but are understandable at this point as well.

First a router instance has to be constructed, to which the route tree needs to be provided. Additionally, the router context has to be configured with initial values. This has to include everything that routes need to have access to before even rendering their component. In my case this contained TanStack query's query client, which is constructed here, the auth context, and a default callback that returns the given routes title. The latter will be overridden in each route and will be used for breadcrumb generation. The auth context is undefined at this point and will be updated later. Here I also set up a default view transition that provides a smoother animation between route switching and scroll restoration. Moreover, I configured the loader stale time to 0 (never stale) as the query client will handle that instead of the router (both share this concept). Finally, I changed the default preload method to 'intent', which has the effect of launching loader preloads after 50 milliseconds when a user hovers on the link of the destination route.

Next a component I called `InnerApp` has to be created. This component's only purpose is to have access to the `useAuth` hook and return the router provider with the router instance and a now existing valid authentication context. The router provider will take care of rendering child routes.

Moving on, I created an `App` component which wraps the previous `InnerApp` with the auth provider, the theme provider, and query client provider. The query client provider has access to the same query client as the router. Also, I included the `Toaster` component from the `Sonner` package created by Emil Kowalski alongside the `InnerApp`. This package

an easy-to-use and opinionated way of making toast notifications. I primarily use it to provide feedback to the user from background tasks such as API calls.

Last but not least the App itself is surrounded with React's strict mode component that enables development only extra behaviours with the aim of finding common bugs. In production, it has no effect and is turned off. At the end of all this, the component chain gets rendered by the root React DOM object, which finishes the setup process, as the index.html by default includes the main.tsx script.

One last thing I want to talk about here is the root route. It must be named as `__root.tsx` and is the root of the route tree. Here the library's Outlet component has to be returned which will take care of rendering all child routes. Additionally, I implemented application wide, loading, not found and error boundary components to provide a better UX for these scenarios.

For example Figure 14 the not found component that renders when a user navigates to a non-existent route.

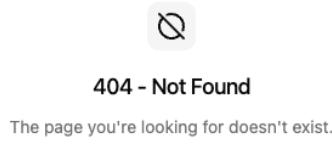


Figure 14: Not found component for missing routes

4.2.2 API client and context providers

After the setup process, I will move on to describing the features mainly in the context, providers, hooks, and lib folders.

Generating the OpenAPI schema of the backend has the great benefit of allowing me to take advantage of the many available libraries specializing in API client code generation. I

investigated a couple of existing solutions, namely Orval, Kubb, and OpenAPI-TS. Orval and Kubb both feature an extendible approach where you can pick what you want to use. However, they both proved to be way more complicated than what I needed, so I opted to use OpenAPI-TS. OpenAPI-TS provides many smaller libraries that provide client-side code generation from OpenAPI schemas. For my use case OpenAPI-React-Query was the one I needed and ended up using. It is a tiny, around 1kb wrapper around TanStack Query and incorporates two other libraries by the same author, OpenAPI-Fetch and OpenAPI-TypeScript. It validates the available API routes inside the query calls, checks parameters and bodies, and can be used to generate TypeScript types for the requests and responses.

Next I generated the API client with the library in the `lib/api` folder and also customized the fetch client used by it. Namely, I created a custom `APIError` class that handles error messages from the backend, configured the base URL from environment variables and the default headers. Additionally, since the backend uses CORS and CSRF protection it requires requests to have a valid XRSF token. This can be obtained by first making a request to the `/sanctum/csrf-cookie` endpoint. This was also configured in the fetch client to be handled automatically and re-fetch when the token expires.

Moving on, I implemented the custom theme and authentication providers. I was guided by the ShadCN documentation [33] on how to configure light, dark, and system themes using TailwindCSS's functionality. The provider stores the selected theme in and recovers it from the browser's local storage on subsequent visits. Toggling the themes are possible from the account page, and the app defaults to the system theme. Accessing the theme in other components is possible through a custom `useTheme` hook.

Afterwards I created the authentication provider for auth state management. It provides through the `useAuth` hook methods for logging in, registering, logging out, re-fetching user context and deleting the account. All these methods use the API client wrapping TanStack Query's mutations and query functions. The context encompasses two React states, one a boolean for quick checks, and the other containing the logged-in user's information. I also store the state in the browser's local storage and recovered upon mounting by an `useEffect` call so that the user remains logged in on revisits.

4.2.3 Authentication, dashboard and account routes

In this block I will mainly be presenting the authentication and account related pages that I created on the frontend. To begin with, a use case diagram on Figure 15 shows in detail the authentication capabilities.

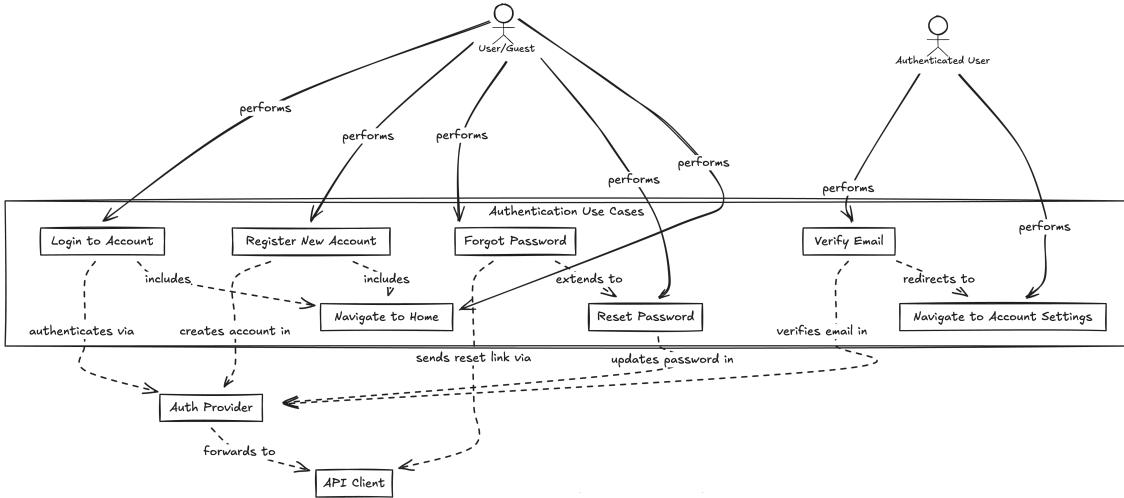


Figure 15: Authentication use cases

Guests or unauthenticated users can log in, register and ask for a forgot password link which redirects them to a form allowing them to reset their password. Authenticated users can verify their email's which they automatically get after registration or can request a new one on the account page, which I will talk about later. Obviously if the user clicks on the verification link in their email while being unauthenticated they will first be redirected to the login page.

When a user first visits the frontend a simple index page greets them with directions to login, or register, or go to the dashboard if already logged in. The index page as shown for an unauthenticated user can be seen on Figure 16.

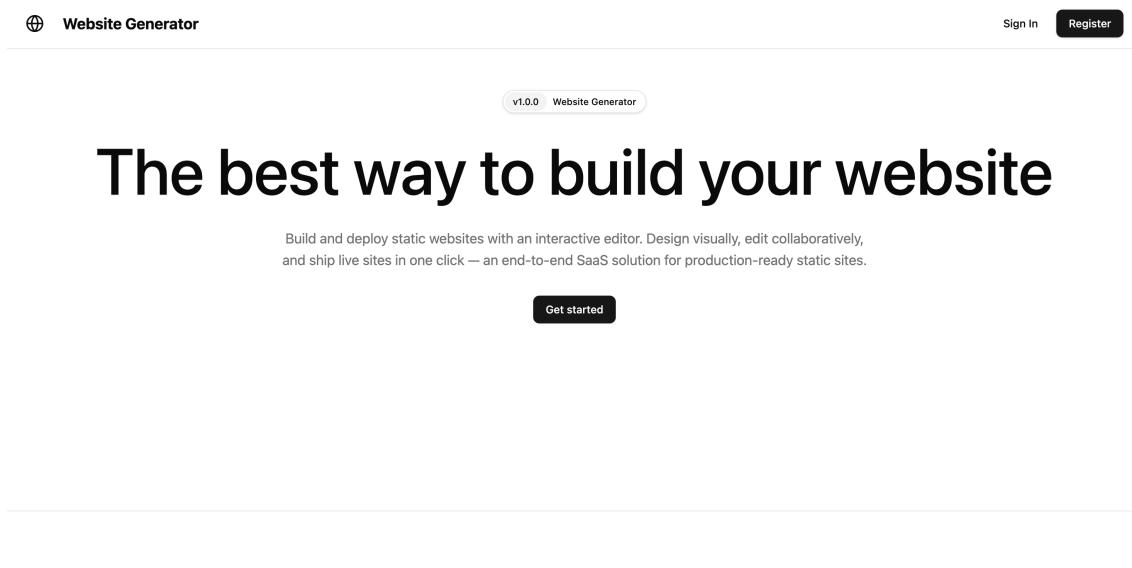


Figure 16: Index page for an unauthenticated user

From the landing page, the user can directly log in and or register. As stated previously I used React Hook Form with Zod for creating all my forms. This ensures that there is proper client side validation, and can be used to display meaningful error messages to the

user. All forms also employ a loading state to prevent multiple submissions while a previous request is running. I also display success and error responses using the aforementioned toast notifications in the bottom right corner of the screen. The registration page with an example for validation errors can be seen on Figure 17.

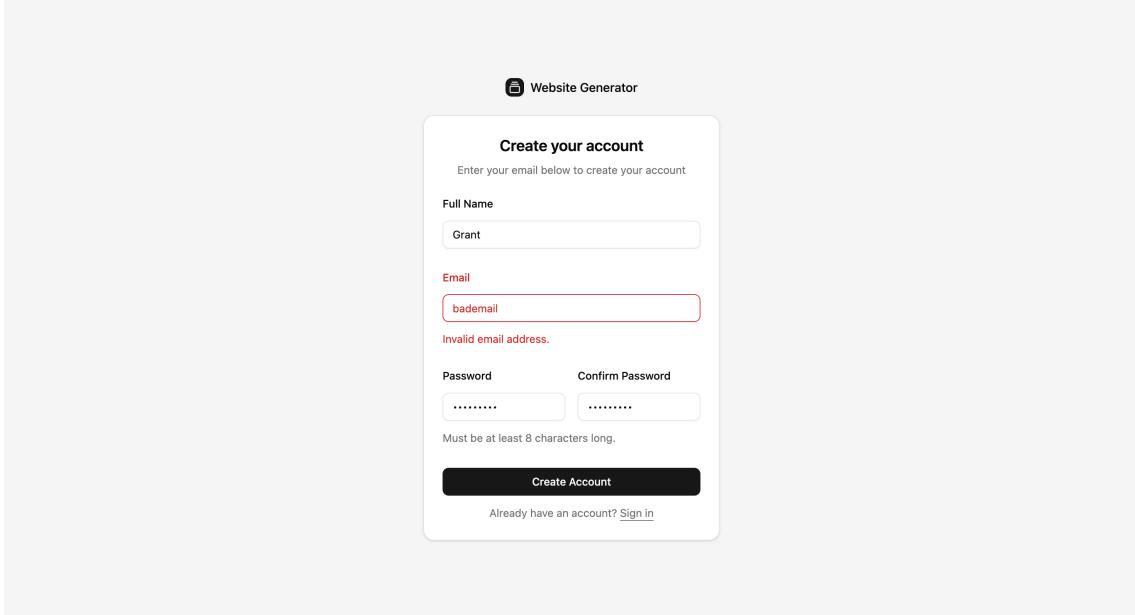


Figure 17: Register page with validation errors

The rest of the forms look almost identical to the registration one, so I won't show each one-by-one unnecessarily. One thing worth mentioning is that both the registration and forgot password links send out email. Such an example email after registration captured by Mailpit can be seen on Figure 18. On the sidebar the incoming password reset links can be seen as well.

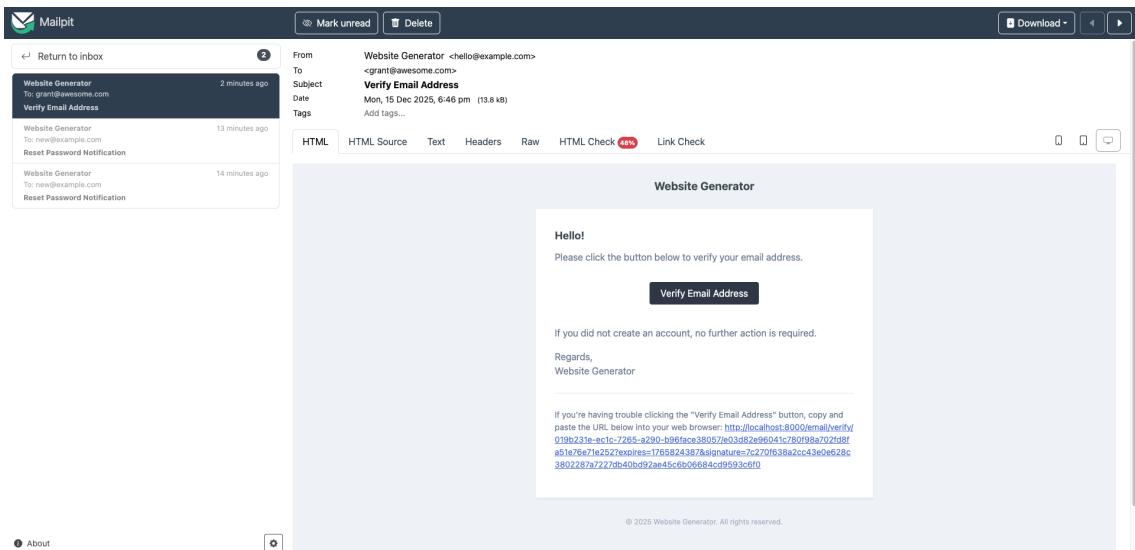


Figure 18: Registration email seen in Mailpit

After logging in, the user can access the main application routes grouped by the `_app` folder. They are all surrounded with a layout component that has two main responsibil-

ties. First check if the user has successfully logged in, if not redirect them to the login route. Secondly this is where the title callback I mentioned in Subsection 4.2.1 is used. The layout component collects the result of all title callbacks from the router state, which include the title names for all rendered child routes. Afterwards it dynamically builds the breadcrumb trail displayed to the user.

The first such route is the dashboard, which displays a few statics and quick actions. Figure 19 showcases it in detail.

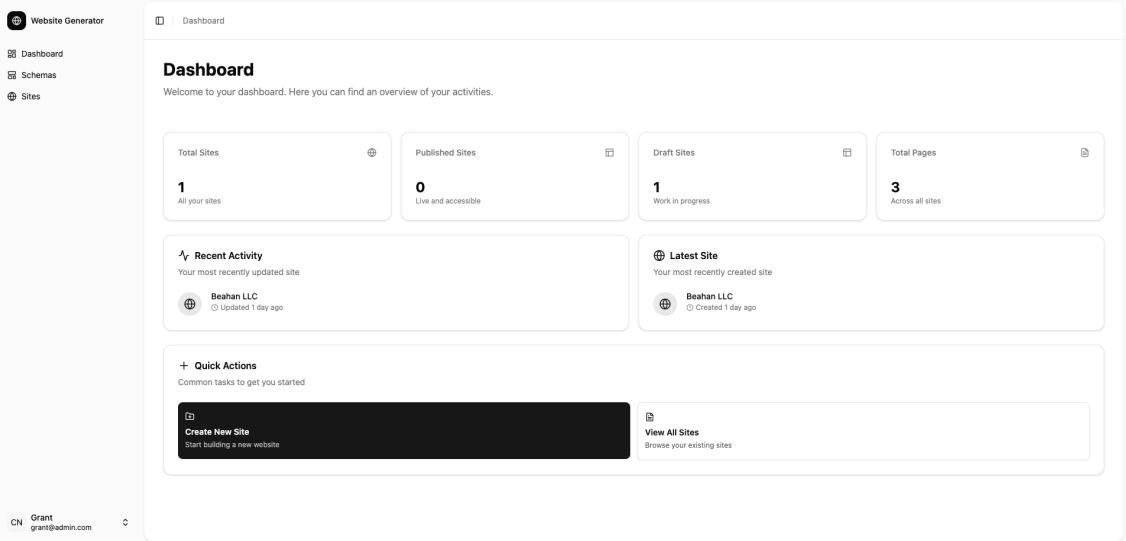


Figure 19: The dashboard page

Next let me provide an overview of account settings, which can be accessed via a pop-over by clicking on the user section of the sidebar. Another use case diagram showcasing the account related options can be seen on Figure 20.

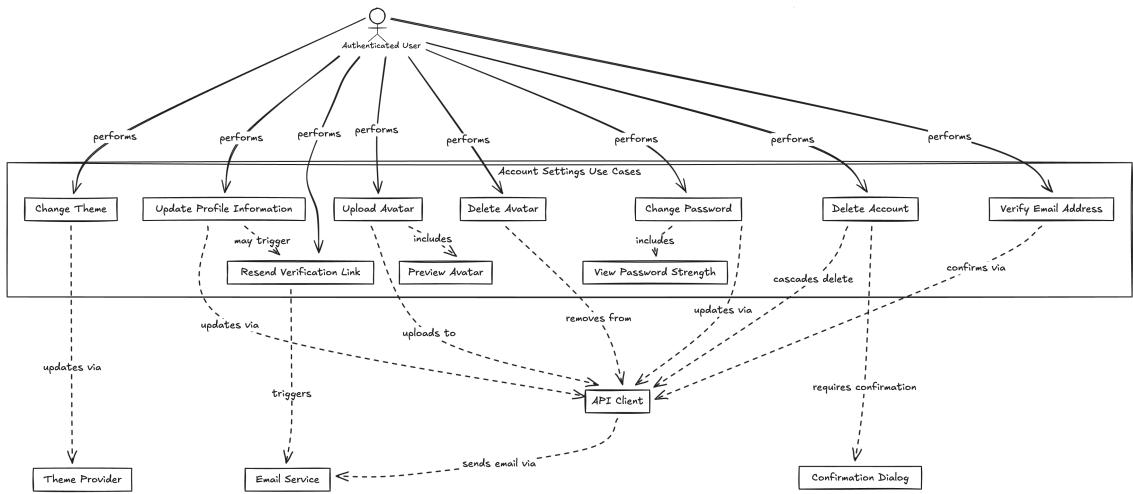


Figure 20: Account settings use cases

The user is able to switch the application between light, dark, and system themes. They can also update their profile information which includes changing their display name and email. Changing the email automatically sends a verification email for the new address. They can upload an avatar via a drag and drop field or the system file chooser, which can

be previewed before saving. A tab which show the verification status is available, and if unverified a new link can be sent from here. Password's can be changes using a form, that shows the new password's strength using simple metrics like length and contained special characters. After changing it the user will be logged out automagically and prompted to log in again. Lastly the account can be deleted if the user wishes so. This will first show a confirmation dialog to avoid accidental deletion.

Two pictures comparing light and dark modes can be seen on Figure 21. The light mode picture on Subfigure 21a showcases the avatar upload tab, while the dark mode picture on Subfigure 21b presents the password change form. The light mode picture also displays, how to get to the account settings page.

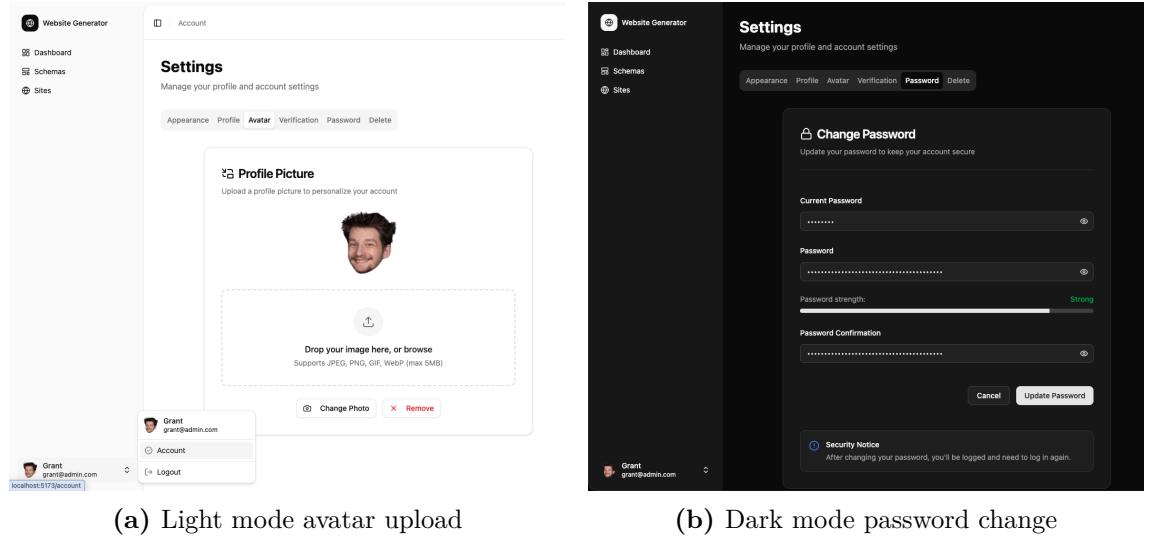


Figure 21: Dark and light modes

4.2.4 Schemas, sites, and pages

4.2.5 Editor

4.3 Proxy application

Chapter 5

Testing

5.1 Backend testing

5.2 Frontend testing

5.3 Usage example

Chapter 6

Production considerations

Chapter 7

Summary

Acknowledgements

Bibliography

- [1] ISO/IEC22123-1. Information technology — Cloud computing Part 1: Vocabulary. Standard, International Organization for Standardization, Geneva, CH, February 2023.
- [2] Peter Mell, Tim Grance, et al. The nist definition of cloud computing. 2011.
- [3] Saiqa Aleem, Rabia Batool, Faheem Ahmed, Asad Khatak, and Raja Muhammad Ubaid Ullah. Architecture guidelines for saas development process. In *Proceedings of the 2017 International Conference on Cloud and Big Data Computing*, pages 94–99, 2017.
- [4] Javier Espadas, David Concha, and Arturo Molina. Application development over software-as-a-service platforms. In *2008 the third international conference on software engineering advances*, pages 97–104. IEEE, 2008.
- [5] Rouven Krebs, Christof Momm, and Samuel Kounev. Architectural concerns in multi-tenant saas applications. *Closer*, 12:426–431, 2012.
- [6] WeiTek Tsai, XiaoYing Bai, and Yu Huang. Software-as-a-service (saas): perspectives and challenges. *Science China Information Sciences*, 57(5):1–15, Mar 2014. DOI: <https://doi.org/10.1007/s11432-013-5050-z>. URL <https://link.springer.com/article/10.1007/s11432-013-5050-z>.
- [7] Farhan Aslam. The benefits and challenges of customization within saas cloud solutions. *American Journal of Data, Information, and Knowledge Management*, 4(1):14–22, 2023.
- [8] Wei-Tek Tsai, Yu Huang, Xiaoying Bai, and Jerry Gao. Scalable architectures for saas. 04 2012. DOI: 10.1109/ISORCW.2012.44.
- [9] Matt Stauffer. *Laravel: Up & running: A framework for building modern php apps*. " O'Reilly Media, Inc.", 2019.
- [10] Zoltán Subecz. Web-development with laravel framework. *Gradus*, 8(1):211–218, 2021.
- [11] Martin Bean. *Laravel 5 essentials*. Packt Publishing Ltd, 2015.
- [12] Laravel 12 documentation. <https://laravel.com/docs/12.x>, 2025. Accessed: 2025-12-08.
- [13] Spatie. <https://spatie.be/>, 2025. Accessed: 2025-12-08.
- [14] Filament php. <https://filamentphp.com/>, 2025. Accessed: 2025-12-08.

- [15] Arshad Javeed. Performance optimization techniques for reactjs. In *2019 IEEE International Conference on Electrical, Computer and Communication Technologies (ICECCT)*, pages 1–5. IEEE, 2019.
- [16] Madhuri A Jadhav, Balkrishna R Sawant, and Anushree Deshmukh. Single page application using angularjs. *International Journal of Computer Science and Information Technologies*, 6(3):2876–2879, 2015.
- [17] Sagar Ganatra. *React Router Quick Start Guide: Routing in React Applications Made Easy*. Packt Publishing Ltd, 2018.
- [18] Merging remix and react routers. <https://remix.run/blog/merging-remix-and-react-router>, 2024. Accessed: 2025-12-08.
- [19] Tanstack router documentation. <https://tanstack.com/router/latest/docs/framework/react/overview>, 2025. Accessed: 2025-12-08.
- [20] Tanstack query documentation. <https://tanstack.com/query/latest/docs/framework/react/overview>, 2025. Accessed: 2025-12-08.
- [21] Tashi Garg. React query: Revolutionizing data fetching and modernizing user interfaces. *Authorea Preprints*, 2025.
- [22] Mark Nottingham. Rfc 5861: Http cache-control extensions for stale content. <https://datatracker.ietf.org/doc/html/rfc5861>, 2025. Accessed: 2025-12-08.
- [23] Dave Micheal. Leveraging stale-while-revalidate in react query for optimal ux. 2025.
- [24] Mahesh Kumar Bagwani, Virendra Kumar Tiwari, and Ashish Jain. Implementing grapesjs in educational platforms for web development training on aws. *Int. J. Sci. Res. in Multidisciplinary Studies Vol*, 10(8), 2024.
- [25] Craftjs documentation. <https://craft.js.org/docs/overview>, 2025. Accessed: 2025-12-08.
- [26] Arun Gupta Alan Hohn. Getting started with kubernetes. *Kubernetes*. Retrieved May 2019., 24:2019, 2019.
- [27] Marko Luksa. *Kubernetes in action*. Simon and Schuster, 2017.
- [28] Ingress. Kubernetes ingress. <https://kubernetes.io/docs/concepts/services-networking/ingress/>, Nov 2025. Accessed: 2025-12-08.
- [29] Upendra Kanuru and Upendra Kumar Gurugubelli. Modern approach to kubernetes traffic management: Migrating from ingress to gatewayapi. *International Journal of Emerging Trends in Computer Science and Information Technology*, 6(3):1–11, Jul. 2025. DOI: 10.63282/3050-9246.IJETCSIT-V6I3P101. URL <https://www.ijetcsit.org/index.php/ijetcsit/article/view/282>.
- [30] Rahul Sharma and Akshay Mathur. *Traefik API Gateway for Microservices*. Springer, 2020.
- [31] shadcn.io. shadcn.io. <https://www.shadcn.io/>, 2025.
- [32] Bulletproof react. <https://github.com/alan2207/bulletproof-react>, 2025. Accessed: 2025-12-08.
- [33] shadcn. Introduction. <https://ui.shadcn.com/docs>, 2025.